

Accessing Object Representations

Document #: D1839R2
Date: 2019-11-13
Project: Programming Language C++
Core Working Group
Reply-to: Krystian Stasiowski
<sdkrystian@gmail.com>

Contents

1 Abstract	2
2 Revisions	2
2.1 Changes since [P1839R0]	2
2.2 Changes since [P1839R1]	2
3 Motivation	2
4 Problem	2
5 Changes	3
5.1 Examples	3
6 Design Choices	4
6.1 Object contiguity	4
6.2 Preserving <code>reinterpret_cast</code> and <code>static_cast</code> equivalence	4
6.3 The <code>std::launder</code> issue	4
6.4 “Self-representing” objects	4
6.5 Value and access of elements of object representations	5
7 Wording	5
7.1 Memory and object model [intro.object], [intro.memory], [basic.life]	5
7.2 Access to object representations via <code>reinterpret_cast</code> [expr.reinterpret.cast]	6
7.3 Pointer arithmetic [expr.add]	6
7.4 <code>std::launder</code> [ptr.launder]	7
8 Acknowledgements	7
9 References	7

1 Abstract

Allow access to the object representation of an object.

2 Revisions

2.1 Changes since [P1839R0]

- Allow pointer arithmetic on expressions of type `unsigned char*`, `char*` and `std::byte*` when pointing to objects of different type.
- Removed exclusion of the object representation of objects of zero size from appearing in the object representation of their containing object.
- Added multi-dimensional arrays of contiguous-layout types to the definition of contiguous-layout types.
- Slight change to the behavior of `std::launder` for when there are multiple viable objects.

2.2 Changes since [P1839R1]

- P1839R1 approved by EWG.
- Removed contiguous-layout types from wording, this should be tackled by [P1945R0].

3 Motivation

This proposal does not intend to introduce anything new, rather to standardize a common existing practice. Accessing the underlying bytes of an object has been a long-standing practice in C and C++ alike, but, in C++, doing so is typically undefined behavior. With current wording, it is impossible to obtain a pointer to an element of the object representation; an expression such as `reinterpret_cast<char*>(&a)` typically yields a pointer to the original object, and only the type of the expression is changed. This does not represent the intent of CWG, as exemplified by [CWG1314], in which it is stated that access to the object representation is intended to be well-defined.

This has only recently become undefined behavior as of C++17, when [P0137R1] was accepted. This proposal includes a change to how pointers work, notably that they point to objects, rather than just representing an address, and it seems that the proposal neglected to add any provisions to allow access to the object representation of an object.

4 Problem

This issue exists for two primary reasons: casting and pointer arithmetic. Given the following code:

```
int a = 420;
char b = *reinterpret_cast<char*>(&a);
```

The pointer does not bind to any `char` object or element of the object representation. This particular `reinterpret_cast` is exactly equivalent to `static_cast<char*>(static_cast<void*>(&a))` as per [expr.reinterpret.cast] p7 and as such, [expr.static.cast] p13 dictates that the value of the pointer is unchanged and therefore it points to the original object. When the lvalue-to-rvalue conversion is applied to the initializer expression of `b`, the behavior is undefined as per [expr.pre] p4 because the result of such a conversion would be the value of the `int` object (420), which is not a value representable by `char`.

Additionally, if such wording did exist, an object representation as defined by [basic.types] p4 is a sequence of `unsigned char` objects, not an array, and is unsuitable for pointer arithmetic given the current object model.

5 Changes

- Change object representations to be considered an array if the type of the object they represent is contiguous.
 - Objects of type `unsigned char`, `char` and `std::byte` and arrays of such types suffice as being their own object representation to prevent an infinitely recurring property.
 - The value of the elements of an object representation of a type other than `unsigned char`, `char` and `std::byte` is unspecified, otherwise the value of the element is the value of the object they represent.
- Allow a pointer to an object representation to be obtained through a `reinterpret_cast` to `unsigned char`, `char` and `std::byte`.
- Allow a pointer to an object representation to be cast back to a pointer to its respective object via `reinterpret_cast`.
- Specify that `std::launder` will prefer to return a pointer to an object that is not an element of an object representation.
- Allow pointer arithmetic to be performed on pointers to elements of an object representation if the type of the expression is `unsigned char*`, `char*` or `std::byte*`.

5.1 Examples

Here is an example demonstrating the difference:

Before	After
<pre>using T = unsigned char*; int a = 0; T b = reinterpret_cast<T>(&a); // Pointer value unchanged, still // points to the int object T c = ++b; // UB, expression type differs // from element type</pre>	<pre>using T = unsigned char*; int a = 0; T b = reinterpret_cast<T>(&a); // Pointer now points to the first unsigned // char element of the object representation T c = ++b; // This is now a pointer to the second // element of the object representation ++(*c); // OK</pre>

Another example for arrays:

Before	After
<pre>using T = unsigned char*; int a[5]{}; T b = reinterpret_cast<T>(&a); // Pointer value unchanged, still // points to the array object for (int i = 0; i < sizeof(int) * 5; ++i) b[i] = 0; // UB, expression type differs // from element type</pre>	<pre>using T = unsigned char*; int a[5]{}; T b = reinterpret_cast<T>(&a); // Pointer now points to the first // unsigned char element of the // object representation of the array for (int i = 0; i < sizeof(int) * 5; ++i) b[i] = 0; // OK</pre>

6 Design Choices

6.1 Object contiguity

For pointer arithmetic to work with a pointer to an element of the object representation, it is necessary that the object representation be an array. However, since not all objects are guaranteed to occupy contiguous bytes of storage, the object representation may only be an array if the corresponding object occupies contiguous bytes of storage. It would be useful to expand the guarantee for which objects occupy contiguous storage, and therefore the subset of objects that can have their object representation accessed, however, this will be addressed in later paper [P1945R0].

6.2 Preserving `reinterpret_cast` and `static_cast` equivalence

The proposed wording does not preserve the current `reinterpret_cast` and `static_cast` equivalence for object pointer types when `reinterpret_cast` is used to cast a pointer to `unsigned char*`, `char*`, and `std::byte*`. In these cases, `reinterpret_cast` will no longer follow the pointer-interconvertibility rules used by `static_cast`, instead resulting in a pointer to the object representation.

For example:

```
struct S
{
    unsigned char a;
};

void f()
{
    S b{0};
    unsigned char* c = reinterpret_cast<unsigned char*>(&b);
}
```

The current approach taken by this proposal is to only allow for `reinterpret_cast` to convert `&b` into a pointer that points to an element of the object representation of `b`. Preserving the equivalence leads to the question of whether the cast should follow the pointer-interconvertibility rules and result in a pointer to `b.a`, or result in a pointer to the first element of the object representation of `b`. However, breaking the equivalence will prevent `std::memcpy` from being implemented by the user with standard C++, as casting from a `void*` to `unsigned char*` will not result in a pointer pointing to an element of the object representation. It is unclear which of these it should be and is best to be decided on by CWG.

6.3 The `std::launder` issue

Multiple objects may occupy the same storage, in which case the objects' respective object representations will overlap. This presents the issue of deciding to which object `std::launder` returns a pointer. This proposal remedies the issue by prioritizing objects that are not elements of an object representation, and if no such object is found, then a pointer to an unspecified element of the set of viable objects is returned.

6.4 “Self-representing” objects

Certain objects are suitable to act as their own object representations, such as objects of type `unsigned char`, `char` and `std::byte` and arrays of these types. This is to prevent infinite recursion of objects having object representations, as happens with the current word if read pedantically.

6.5 Value and access of elements of object representations

“Self-representing” elements of an object representation of non-array type are specified to have their own value; all other elements of an object representation have an unspecified value. The reasoning for this is quite obvious, as it would be extremely difficult to specify what the value of each element would be. Access of the elements is intended to be well-defined, and is under the proposed wording, however it is up to CWG whether it should be specified explicitly.

7 Wording

All wording is relative to [N4835](#).

7.1 Memory and object model [intro.object], [intro.memory], [basic.life]

7.1.1 Memory locations

Modifying an element of an object representation concurrently with another memory location that it overlaps should be a data race. This sentence specifies that the element of the object representation is the same memory location.

Changes to [intro.memory] p3 sentence 1

- ³ A *memory location* is either an object of scalar type or a maximal sequence of adjacent bit-fields all having nonzero width, and any overlapping elements of an object representation.

7.1.2 Object representations

The type of the elements of the object representation should have the same cv-qualification as that of the object they represent to prevent accidental modification of the object indirectly. Additionally, if the object occupies contiguous bytes of storage, then we could consider the object representation to be an array, and thereby make pointer arithmetic well-defined. Certain objects are said to represent themselves so that the object representation does not have an object representation of its own. The value of the elements that do not represent themselves is left unspecified, as specifying it would be effectively impossible. Lastly, it is specified that an object representation appears in an enclosing object representation to make it useful for introspection.

Remove [basic.types] p4 sentence 1

- ⁴ The *object representation* of an object of type **T** is the sequence of **N unsigned char** objects taken up by the object of type **T**, where **N** equals **sizeof(T)**.

Insert a new paragraph below [intro.object] p1

- ² The *object representation* of an object **a** of type cv **T** is a sequence of **N cv unsigned char** objects that occupy the same storage as **a**, where **N** is equal to **sizeof(T)**. The sequence is considered to be an array of **N T** if the object occupies contiguous bytes of storage. The object representation of an object of type **unsigned char**, **char**, **std::byte**, or an array of such types (ignoring cv-qualification), is itself. Unless an object representation is of an object of type **unsigned char**, **char** or **std::byte** (ignoring cv-qualification), the value of the elements of the object representation is unspecified. The object representation of an object nested within an object **o** is guaranteed to appear in the object representation of **o**.

7.1.3 Overlapping objects

This ensures that an object representation and its elements may exist concurrently with the object they represent, as they occupy the same storage.

Changes to [intro.object] p9

- 9 [...] Two objects with overlapping lifetimes that are not bit-fields may have the same address if one is nested within the other, or if at least one is a subobject of zero size and they are of different types, or if at least one is an element of an object representation; otherwise, they have distinct addresses and occupy disjoint bytes of storage.

7.1.4 Identical element lifetime

Specifying the lifetime of an object representation explicitly ensures that its lifetime will begin and end with that of its corresponding object meaning it need not be preserved after the object it represents is destroyed. The lifetime does not begin during construction to match the wording of [class.ctor] p2.

Insert a new paragraph below [basic.life] p2

- 3 The lifetime of the elements of the object representation of an object begins when the lifetime of the object begins. For class types, the lifetime of the elements of the object representation ends when the destruction of the object is completed, otherwise, the lifetime ends when the object is destroyed.

7.2 Access to object representations via `reinterpret_cast` [expr.reinterpret.cast]

7.2.1 Obtaining a pointer to the object representation

As the current wording for `reinterpret_cast` will result in an unchanged pointer value when attempting to access an object representation, adding a special case to `reinterpret_cast` is necessary to be able to obtain a pointer to the object representation, and vice versa. However, this should only happen if the object representation of the object we are casting has an object representation on which we can perform pointer arithmetic.

Replace [expr.reinterpret.cast] p7

- 7 An object pointer can be explicitly converted to an object pointer of a different type. When a prvalue `v` of object pointer type is converted to the object pointer type “pointer to `cv T`”, the result is `static_cast<cv T*>(static_cast<cv void*>(v))`.
- 7 A prvalue `v` of object pointer type “pointer to `cv1 T1`” pointing to an object `a` can be explicitly converted to an object pointer of a different type “pointer to `cv2 T2`”, where `cv2` is the same `cv`-qualification as, or greater `cv`-qualification than `cv1`, the result of which is defined as follows:
- (7.1) — If `a` occupies contiguous bytes of storage and `T2` is `unsigned char`, `char` or `std::byte`, the result is a pointer to the first element of the object representation of `a`.
- (7.2) — Otherwise, if `a` points to the object representation of an object `b` of type `T2` (ignoring `cv`-qualification), or the first element thereof, the result is a pointer to `b`.
- (7.3) — Otherwise, the result is `static_cast<cv2 T2*>(static_cast<cv2 void*>(v))`.

7.3 Pointer arithmetic [expr.add]

7.3.1 Differing element and expression type

In order to make accessing an object representation using a type other than `unsigned char` well-defined, it must be allowed for the type of the expression to differ from that of the object pointed to in cases where the type of the pointer is `char*` or `std::byte*`, as the pointer points to an object of type `unsigned char`.

Replace [expr.add] p6

- 6 For addition or subtraction, if the expressions `P` or `Q` have type “pointer to `cv T`”, where `T` and the array element type are not similar, the behavior is undefined.

- ⁶ For addition and subtraction where `P` or `Q` have type “pointer to cv `T`” and point to an object `o`, one of the following shall hold:
- (6.1) — `T` is similar to the type of the `o`, or
 - (6.2) — `T` is similar to `unsigned char`, `char` or `std::byte` and `o` is an element of an object representation.
- Otherwise, the behavior is undefined.

7.4 `std::launder` [`ptr.launder`]

7.4.1 Multiple overlapping objects

Since multiple elements of an object representation may exist in the same storage, it is unclear which one `std::launder` would return if such a situation were to arise. Specifying that it is implementation-defined (or possibly unspecified) gives this defined, but unclear behavior.

Changes to [`ptr.launder`] p3

- ³ *Returns:* A value of type `T*` that points to `X`. If multiple such objects exist, the result is the object in the set of possible objects that is not an element of an object representation. Otherwise, it is implementation-defined to which object in the set the result points.

8 Acknowledgements

Thank you to Jason Cobb, John Iacino, Marcell Kiss, Killian Long, Theodoric Stier, and everyone who participated on the `std`-proposals mailing list for the countless reviews and suggestions. Additionally, I would like to thank Professor Ben Woodard for his grammatical review. Lastly, I’d like to thank Timur Doumler for presenting this paper at Belfast, among other things.

9 References

- [CWG1314] Nikolay Ivchenkov. 2011. Pointer arithmetic within standard-layout objects.
<https://wg21.link/cwg1314>
- [P0137R1] Richard Smith. 2016. Core Issue 1776: Replacement of class objects containing reference members.
<https://wg21.link/p0137r1>
- [P1945R0] Krystian Stasiowski. 2019. Making More Objects Contiguous.
<https://wg21.link/p1945r0>